

Problem 3

When `myattackermalware()` returns, it returns to where the function originally returned, since in `myattacker` the original function return value is set to `myattackermalware()`, and then `myattackermalware()` is set to return to the original return value. The victim process continues to execute as normal, and eventually returns to `main` as intended. If `myattackermalware()` return value is not set to anything in `myattacker`, then it has a return value of `0x0`, and does nothing, causing `XINU` to break and become non-responsive. `Myattackermalware()` pid is 3, the same as the victim pid because it is being run by the victim process.

Problem 4

In my implementation of monitoring CPU usage, when a process is switched in, the value of `clktimefine` is stored in a global variable called `clktimeswitch`. When the process is switched out, the difference between `clktimeswitch` and `clktimefine` is added to `prcpuused`.

When rerunning the test cases for Problem 4.1 from lab1, each process that called `printloop()` took 28-29 milliseconds of cpu time, and completed in the order ABCD. For Problem 4.3 from lab1, each process took between 26-30 milliseconds, and completed in the order CDBA.

Problem 5

I implemented the dynamic priority scheduling by modifying the `insert` function to check if the queue key was less than or equal to key to be inserted, and then changing the 3rd argument in all calls to `insert` from the priority to the `cpuused` time of the process, which was initialized to 1. The Null process is only run when no other processes are ready, and this is done by initializing its `cpu used` time to a large value.

When running `cpu-intensive()`, all processes used the same amount of cpu time (21 milliseconds), which shows the fairness in sharing of the 4 processes. When running `io-intensive()`, the processes all were within 2 milliseconds of each other, a range of 13-15 milliseconds. However the processes did not stay in their initial order, due to the `prcpuused` not being set precisely on the beginning/end of a millisecond. Overall it is still a fair scheduler, scheduling the ones it perceived to have fun less ahead of those which have run longer.

For mixed, 2 `cpu-intensive` and 2 `io-intensive` processes, the dynamic scheduler was fair to both `cpu-intensive` and `io-intensive` individually, `cpu-intensive` both taking up 23 milliseconds, but bouncing back and forth between the two, and `io-intensive` taking up 13 and 14 milliseconds. Sharing between the 2 groups would be considered "fair" in that the `cpu intense` processes were given priority while the `io intense` process were "waiting" for a response, and thus finished in the most efficient and fair way.

Bonus

A solution that mitigates this problem would be to allow the newly created processes to run for a set period of time and then switching it to the next newly created process. This time would be long enough to prevent the long running processes from starving.